

# LangGraph Multi-Agent Coding Exercise: Research Assistant

## Problem Statement

Build a multi-agent research assistant using LangGraph that helps users gather information about companies. The system should have multiple specialized agents working together, support follow-up questions, and request human clarification when queries are ambiguous.

## Requirements

### 1. Agent Architecture

Implement **4 specialized agents**:

#### a) Clarity Agent

- Analyzes if the user's query is clear and specific
- Checks if company name is mentioned or if query is too vague
- **OUTPUT:** Sets `clarity_status` to "clear" or "needs\_clarification"
- **ROUTES TO:** Interrupt (if unclear) OR Research Agent (if clear)

#### b) Research Agent

- Searches for company information (news, financials, recent developments)
- This agent should derive this information from a search tool (Tavily MCP would be preferred) or you can use the mock data below.
- **OUTPUT:** Returns research findings and assigns a `confidence_score` (0-10)
- **ROUTES TO:** Validator Agent (if confidence < 6) OR Synthesis Agent (if confidence  $\geq$  6)

#### c) Validator Agent

- Reviews research quality and completeness
- Checks if information is sufficient to answer the user's question
- **OUTPUT:** Sets `validation_result` to "sufficient" or "insufficient"
- **ROUTES TO:**
  - Research Agent (loop back if insufficient AND attempts < 3)
  - Synthesis Agent (if sufficient OR max attempts reached)

#### d) Synthesis Agent

- Takes research findings and creates a coherent summary
- Formats the response in a user-friendly way
- Maintains context from conversation history
- **ROUTES TO:** END

## 2. Core Features to Implement

### Multi-turn Conversation

- Maintain conversation history across multiple queries
- Each agent should access previous messages for context
- Support follow-up questions like "What about their competitors?" or "Tell me more about the CEO"

### Human-in-the-Loop (Interrupt)

- When Clarity Agent detects an unclear query, interrupt the workflow
- Request clarification from the user (e.g., "Which company are you asking about?")
- Resume processing after receiving clarification

### State Management

- Use a proper state schema that includes

### Conditional Routing

- Appropriate routing across the CLarity, Research & Validation agents

None

## 3. Mock Data (You Can Use)

Since you don't need real API calls, mock your research with sample data like:

Python

```
mock_research = {
```

```
        "Apple Inc.": {
            "recent_news": "Launched Vision Pro, expanding services
revenue",
            "stock_info": "Trading at $195, up 45% YTD",
            "key_developments": "AI integration across product line"
        },
        "Tesla": {
            "recent_news": "Cybertruck deliveries ramping up",
            "stock_info": "Trading at $242, volatile quarter",
            "key_developments": "FSD v12 rollout, energy storage
growth"
        }
    }
```

### 3. Deliverable

A zipped code repo that demonstrates, at minimum

1. A working LangGraph with 4 agents
2. State schema definition with all required fields
3. 3 conditional routing functions properly implemented
4. Feedback loop from Validator back to Research with attempt counter
5. Interrupt mechanism for unclear queries
6. Multi-turn conversation handling with memory
7. At least 2 example conversation turns showing the loop in action
8. Showcase software engineering best practices (classes, repo structure, formatting, etc)
9. Instructions on how to run the agentic system in [Readme.md](#)
10. This is an open ended problem, so if you have any confusion then make reasonable assumptions and state those in the [Readme.md](#)
11. Anything that can be implemented beyond the above would get a higher rating. Please call that out clearly in [Readme.md](#) in a section titled "Beyond Expected Deliverable"